

Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

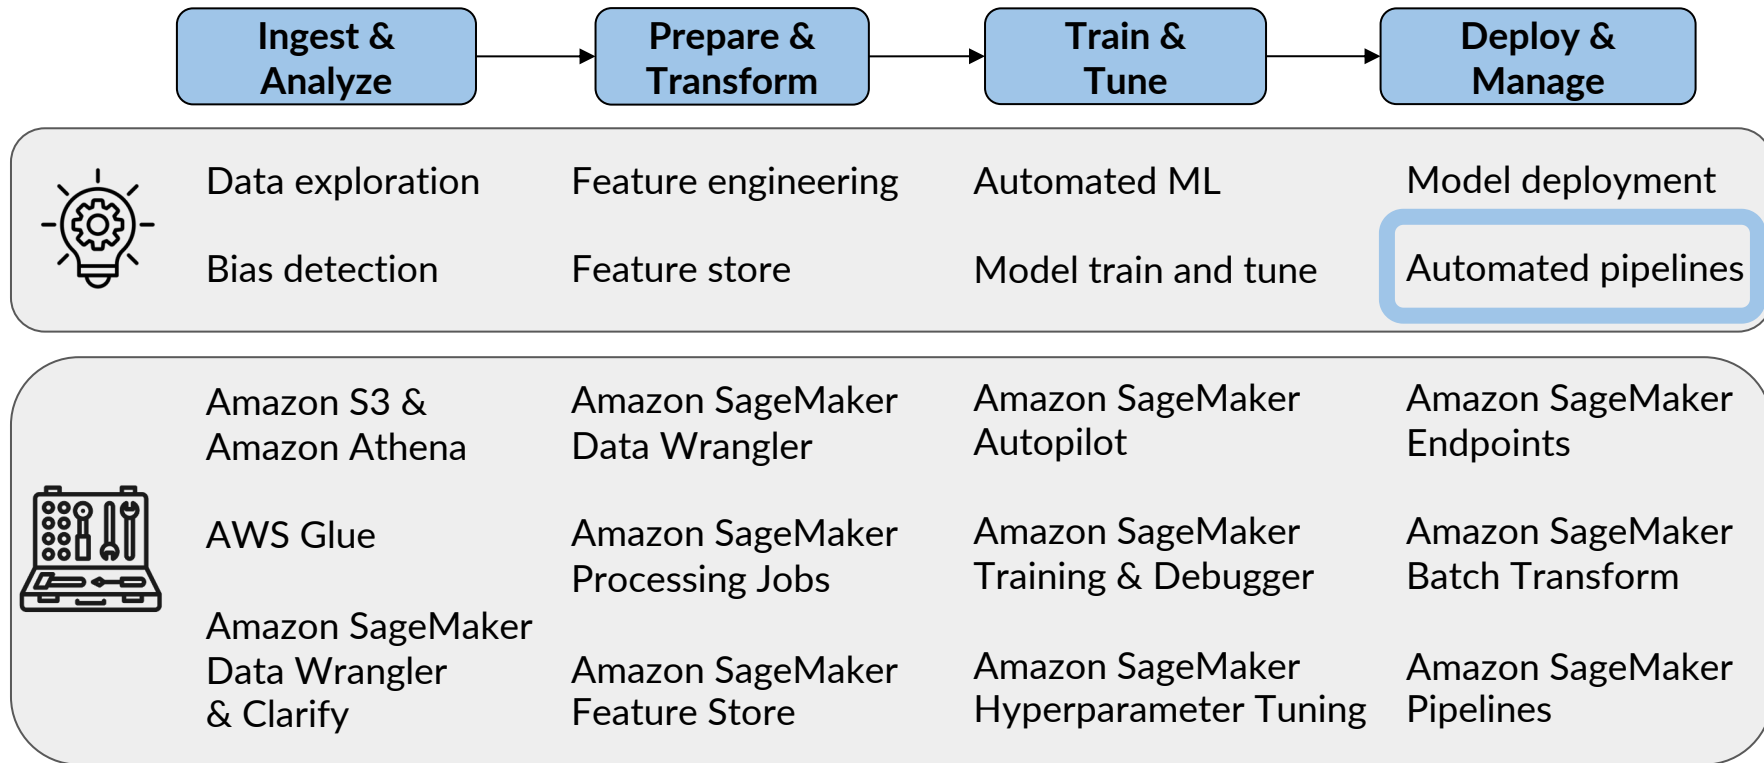
For the rest of the details of the license, see <https://creativecommons.org/licenses/by-sa/2.0/legalcode>

Machine Learning Operations (MLOps)

Overview




Machine Learning Workflow



Path to production for ML models

It's not just **technology**...



Path to production for ML models



Considerations



Machine Learning Development Lifecycle (MLDC)
!= Software Development Lifecycle (SDLC)

Path to production for ML models



Considerations



Machine Learning Development Lifecycle (MLDC)
!= Software Development Lifecycle (SDLC)



A Model may be a small part of an overall solution

Path to production for ML models



Considerations



Machine Learning Development Lifecycle (MLDC)
!= Software Development Lifecycle (SDLC)



A Model may be a small part of an overall solution



Multiple personas spanning the MLDC

Path to production for ML models



Considerations



Machine Learning Development Lifecycle (MLDC)
!= Software Development Lifecycle (SDLC)



A Model may be a small part of an overall solution



Multiple personas spanning the MLDC



Integration with traditional IT practices

Path to production for ML models



Challenges



Culture & Lack of cross functional teams

Path to production for ML models



Challenges



Culture & Lack of cross functional teams



Integration into client applications & existing tooling

Path to production for ML models



Challenges



Culture & Lack of cross functional teams

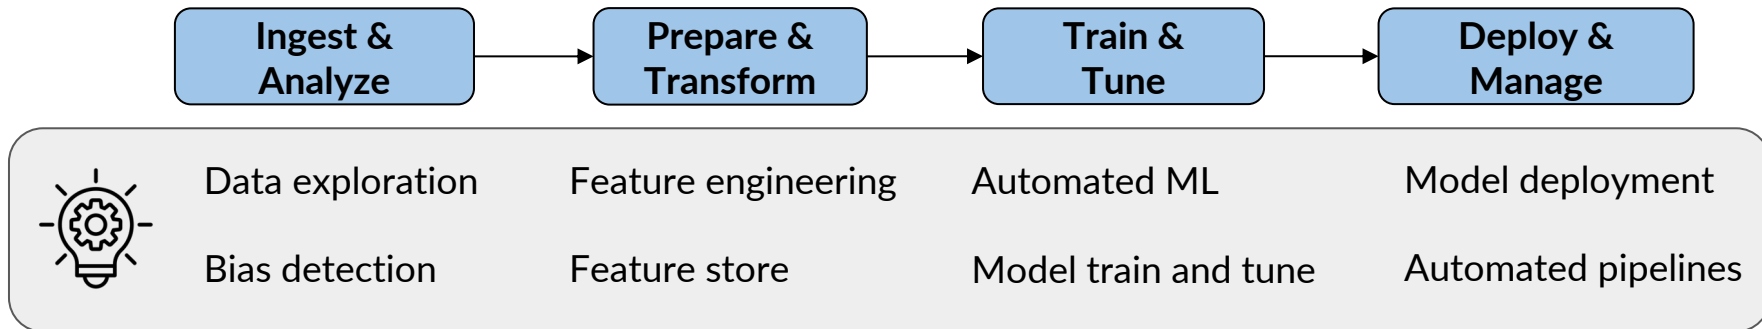


Integration into client applications & existing tooling



Multiple disparate pipelines & dependencies (ex. Code, Data, Training)

Machine Learning Workflow



Operationalizing Machine Learning

Goals

Accelerate the path to production:

- ☐ Reduce manual hand-offs between steps
- ☐ Increase automation within steps
- ☐ Orchestrate the workflow

Operationalizing Machine Learning

Goals

Accelerate the path to production:

- ☐ Reduce manual hand-offs between steps
- ☐ Increase automation within steps
- ☐ Orchestrate the workflow

Improve the quality of deployed models:

- ☐ Implement automated workflows with quality gates

Operationalizing Machine Learning

Goals

Accelerate the path to production:

- ☐ Reduce manual hand-offs between steps
- ☐ Increase automation within steps
- ☐ Orchestrate the workflow

Improve the quality of deployed models:

- ☐ Implement automated workflows with quality gates

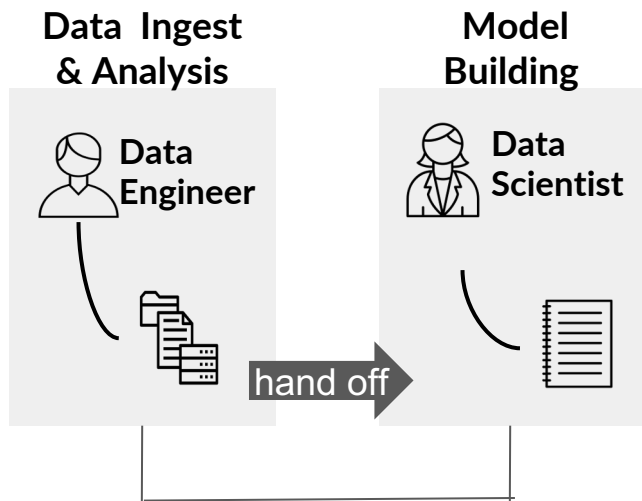
Build resilient, secure, performant, operationally efficient and cost optimized ML solutions

- ☐ Consider aspects unique to ML solutions + Traditional systems engineering considerations

Operationalizing Machine Learning

Path to production

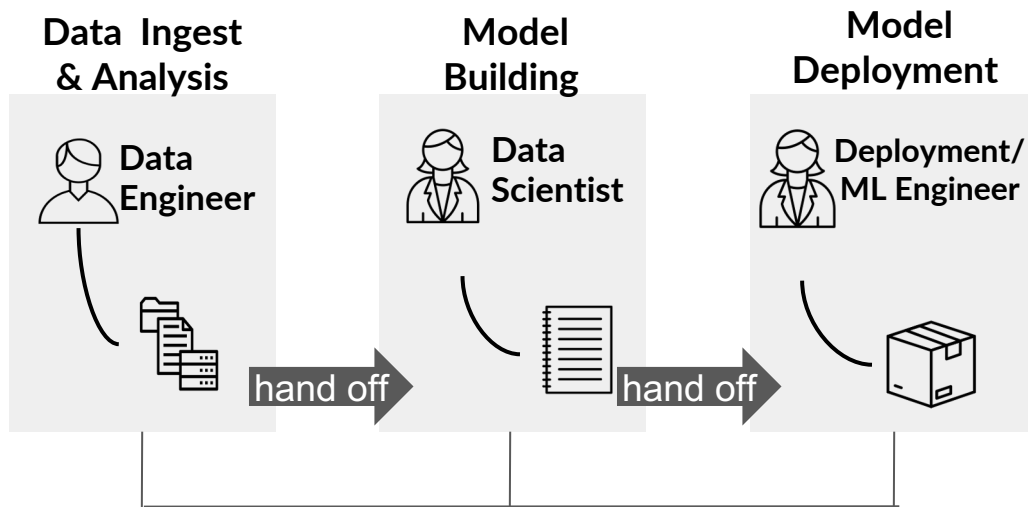
Example workflow with multiple hand offs:



Operationalizing Machine Learning

Path to production

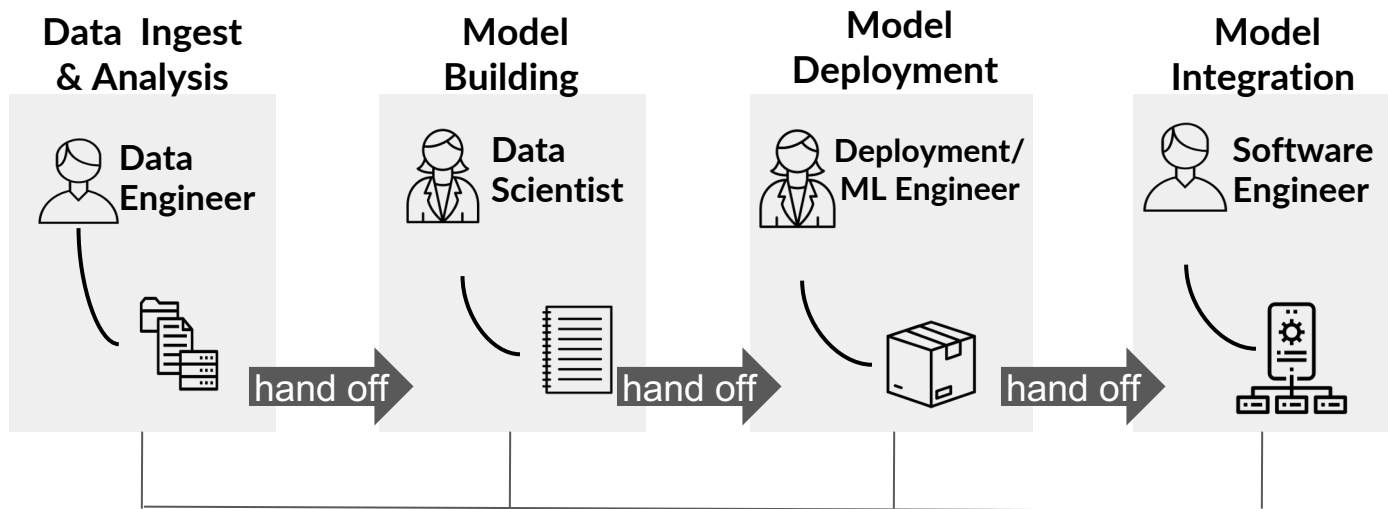
Example workflow with multiple hand offs:



Operationalizing Machine Learning

Path to production

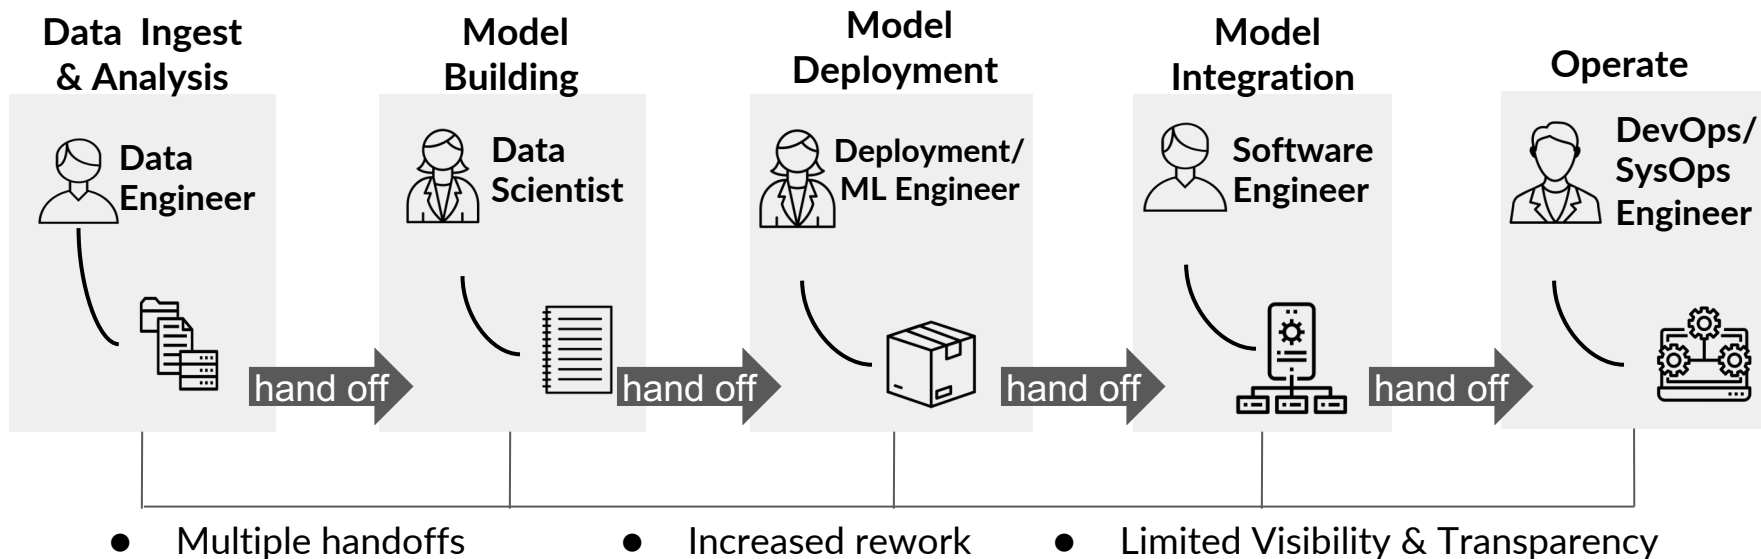
Example workflow with multiple hand offs:



Operationalizing Machine Learning

Path to production

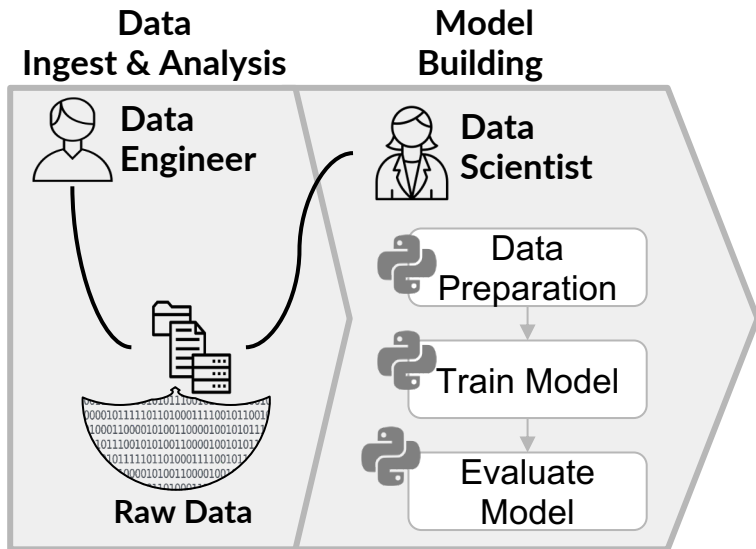
Example workflow with multiple hand offs:



Operationalizing Machine Learning

Accelerate the path to production

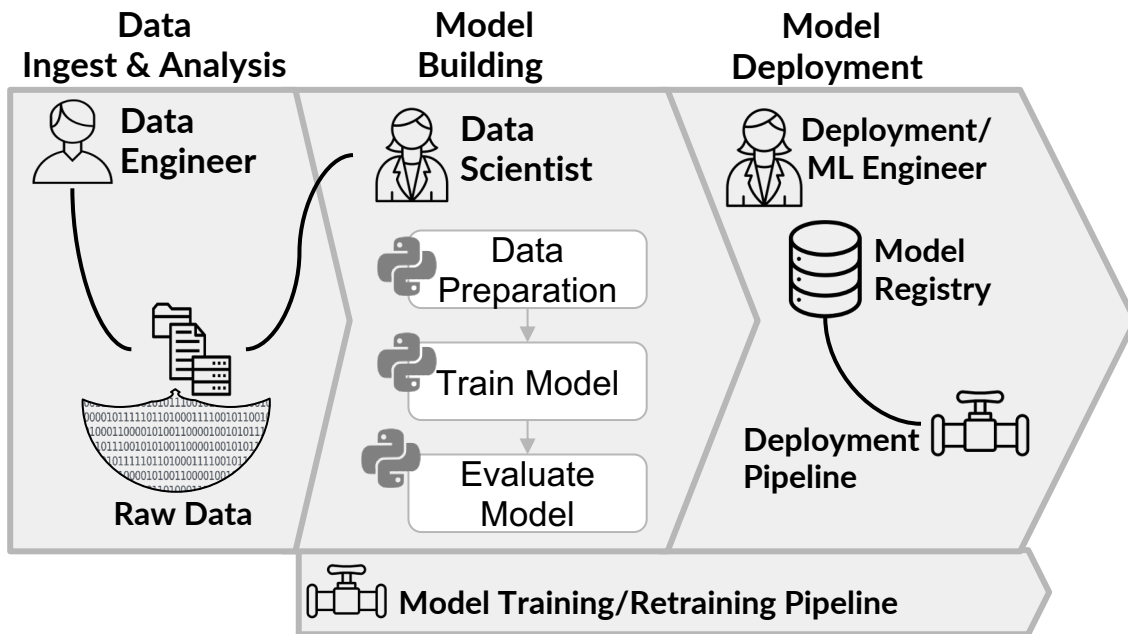
Automating the tasks within a workflow step



Operationalizing Machine Learning

Accelerate the path to production

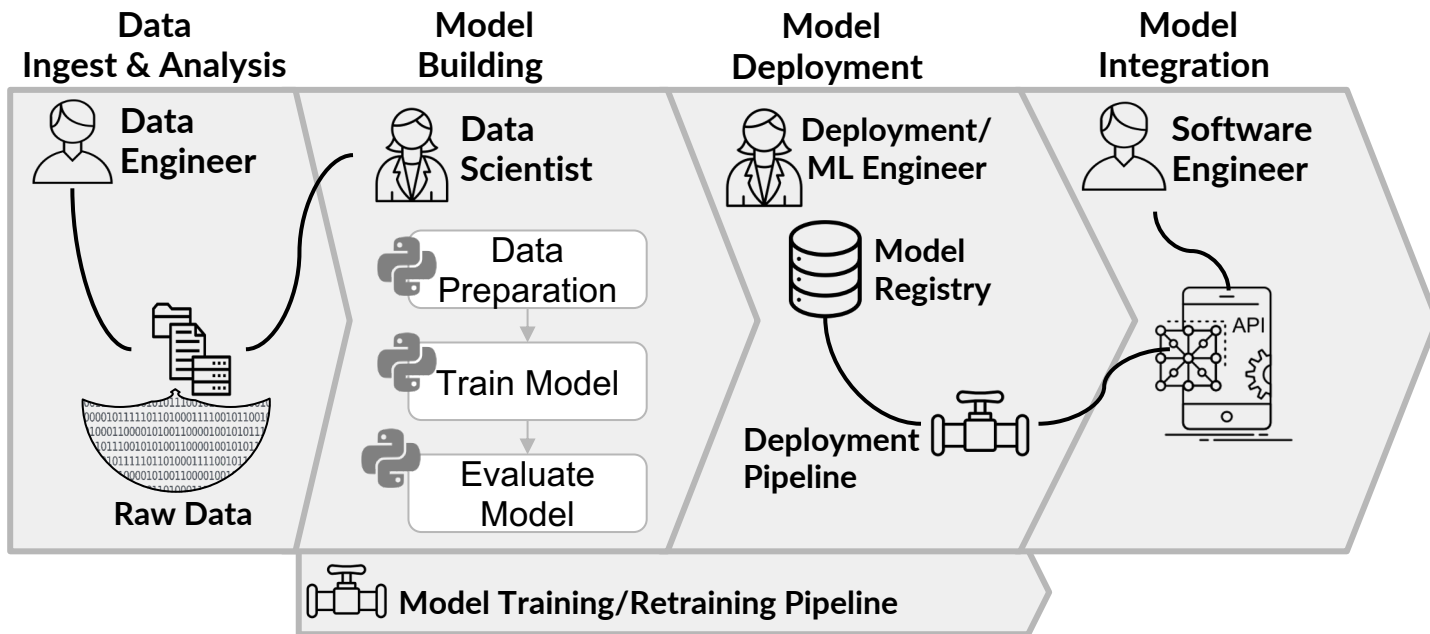
Automating the tasks within a workflow step



Operationalizing Machine Learning

Accelerate the path to production

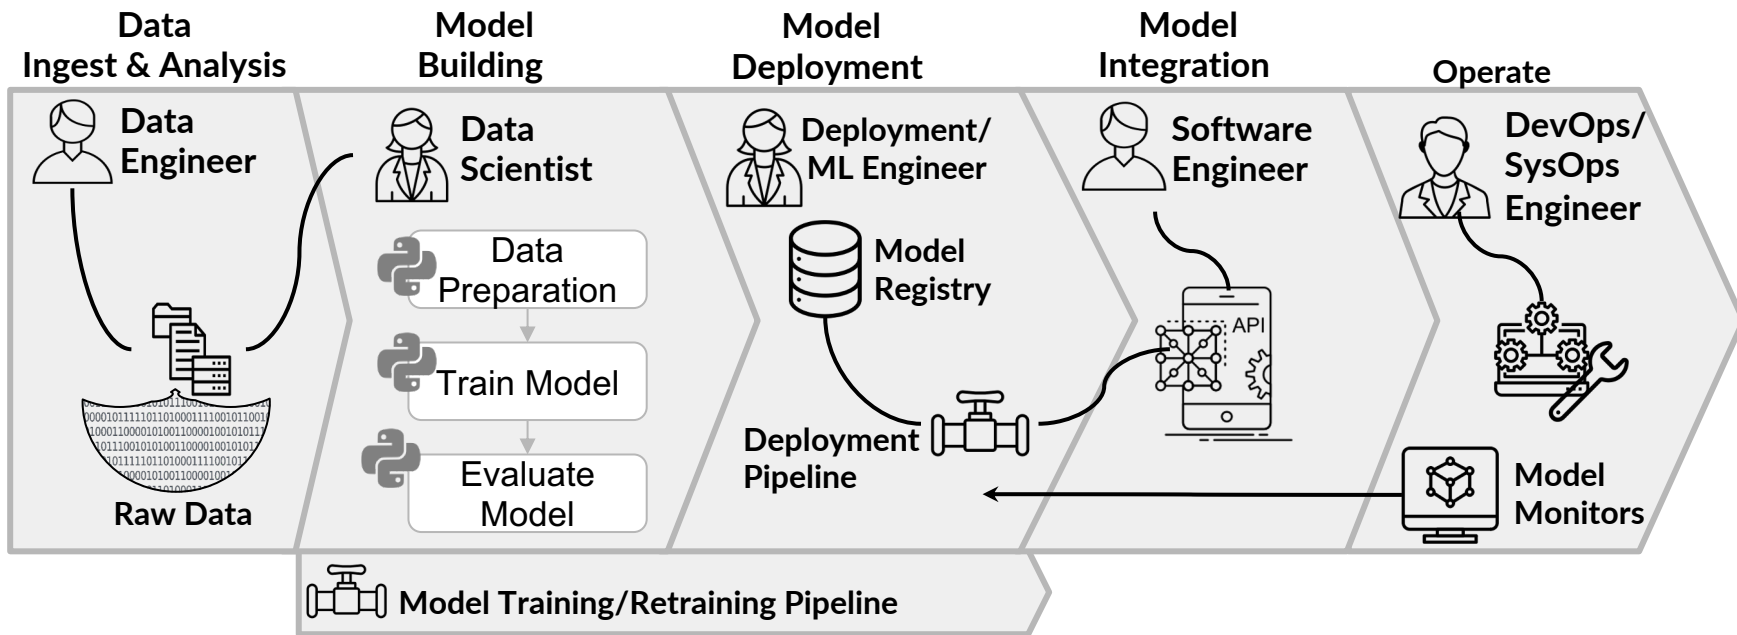
Automating the tasks within a workflow step



Operationalizing Machine Learning

Accelerate the path to production

Automating the tasks within a workflow step




Operationalizing Machine Learning

Accelerate the path to production

Automation vs Orchestration

Automation: Automate a **task** (Ex. Data Preparation) to perform a specific activity or produce defined artifacts based on the inputs or triggers of that task without human intervention



Operationalizing Machine Learning


Accelerate the path to production

Automation vs Orchestration

Automation: Automate a **task** (Ex. Data Preparation) to perform a specific activity or produce defined artifacts based on the inputs or triggers of that task without human intervention

Orchestration: Orchestrate the steps of a workflow that contain a **collection of tasks**


Example: Data Preparation



Operationalizing Machine Learning

Improve the quality of deployed models

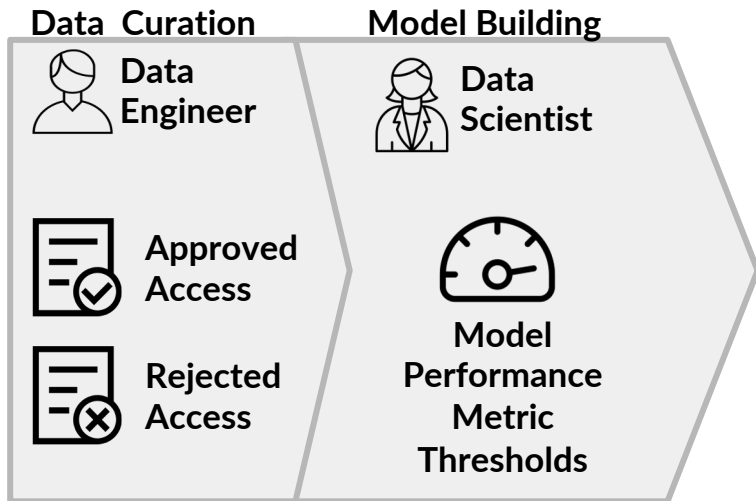
Examples: Automated quality gates



Operationalizing Machine Learning

Improve the quality of deployed models

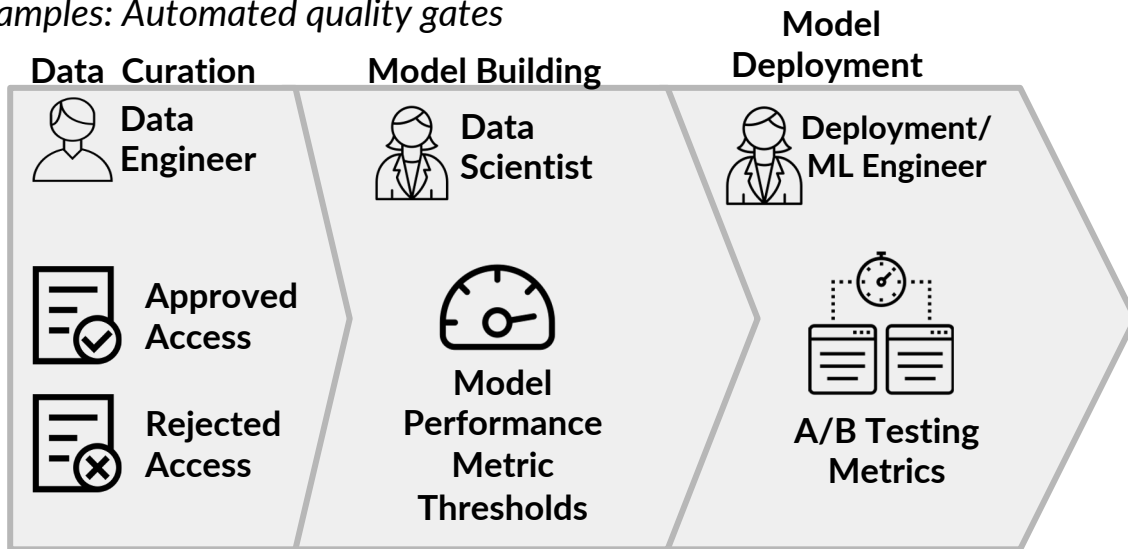
Examples: Automated quality gates



Operationalizing Machine Learning

Improve the quality of deployed models

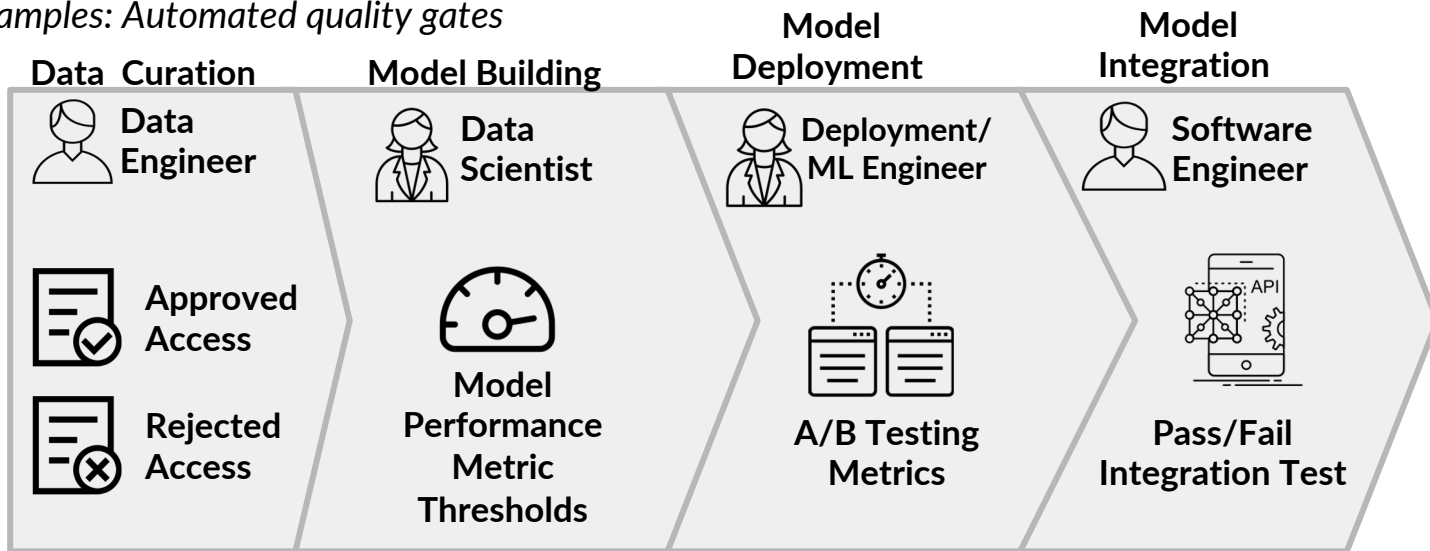
Examples: Automated quality gates



Operationalizing Machine Learning

Improve the quality of deployed models

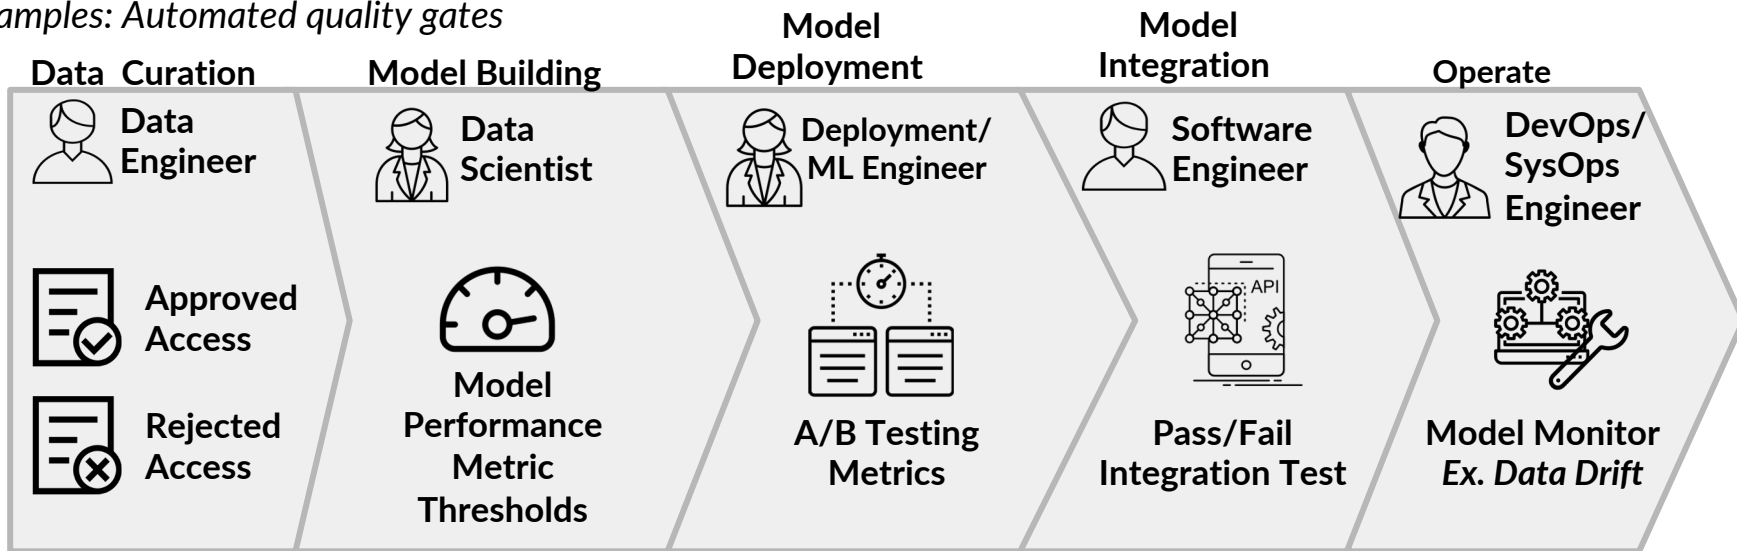
Examples: Automated quality gates



Operationalizing Machine Learning

Improve the quality of deployed models

Examples: Automated quality gates




Operationalizing Machine Learning

Key Considerations




Security




Reliability



Cost
Optimization



Performance
Efficiency



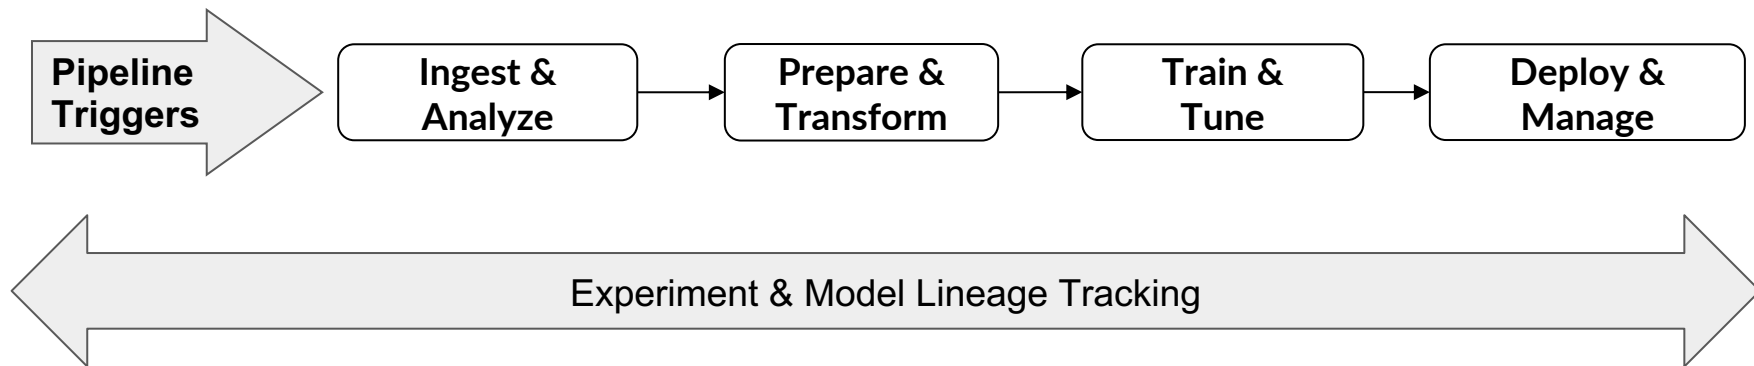
Operational
Excellence

Creating Machine Learning Pipelines



Creating Machine Learning Pipelines

Building Effective Pipelines



Data Tasks

Data Ingestion for Model Development

FROM...

Data
Scientist




Can I get an extract
from our CRM
System?

Data
Engineer





Sure, give me a
couple of days




Data Tasks

Data Ingestion for Model Development






TO...




Data Tasks

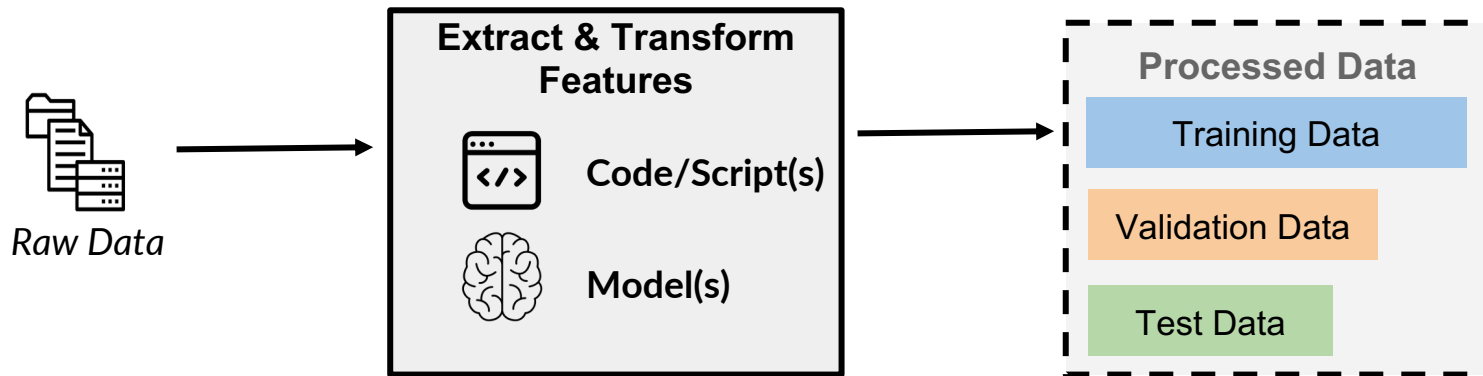
Data Ingestion for Model Retraining




Data Tasks



Data Pre-Processing & Feature Engineering

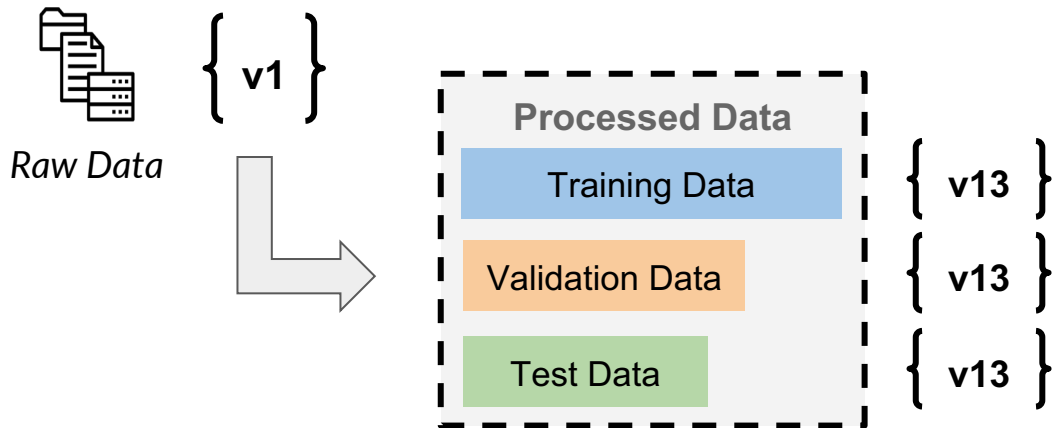


Data Tasks



Data Versioning

Examples:




Data Tasks

Data Validation


Examples:




Data Quality




Statistical Bias



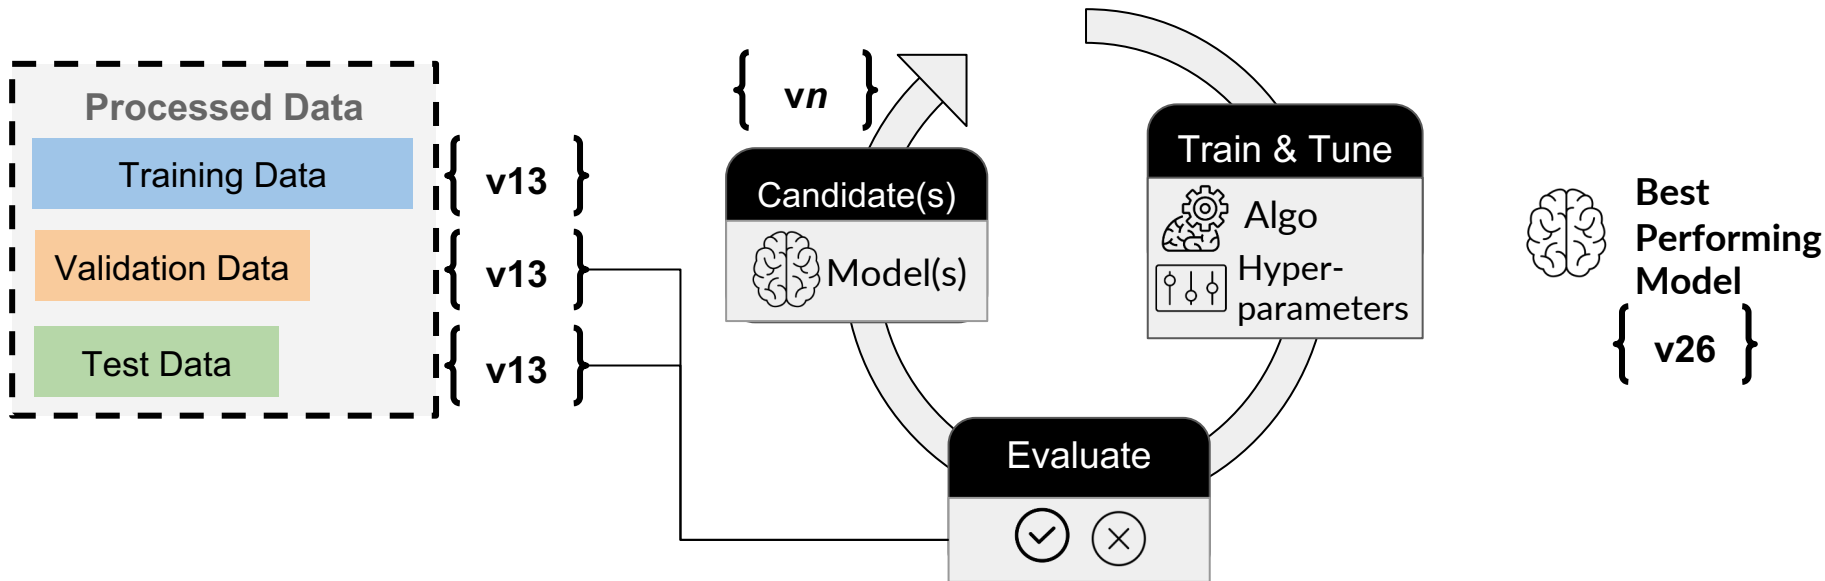
Data Schema




Model Building Tasks



Model Training, Evaluation & Versioning

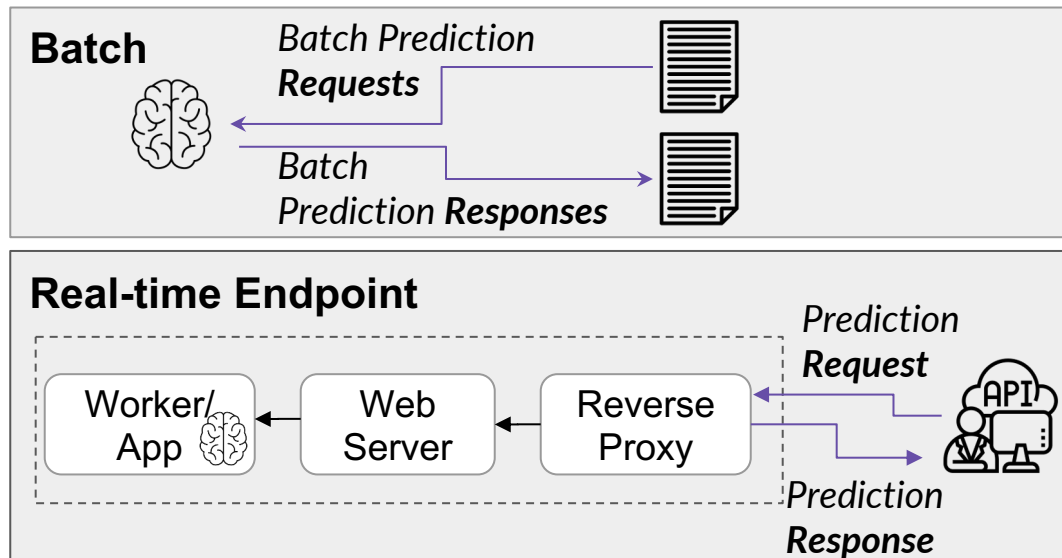


Model Deployment Tasks




Model Deployment & Consumption

Best Performing Model
{ v26 }



Operating Tasks

Logging & Monitoring




Logging →

- Model Data
- System Data




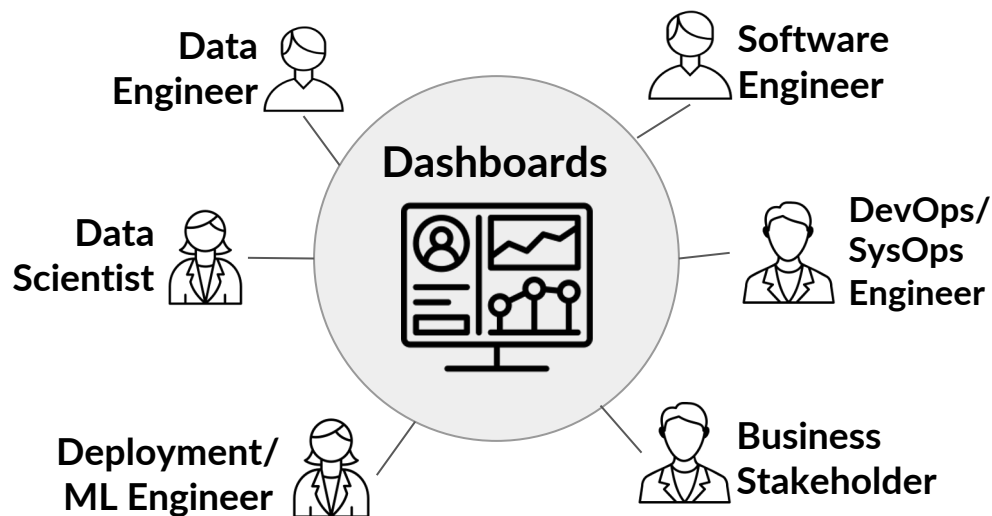
Monitoring →

- Collect Metrics
- Setup Alerts
- Trigger Automated Flows



Operating Tasks

Additional Feedback Mechanisms




- Each persona can have different motivations and needs for monitors, logging and dashboards.
- Examples:
 - Pipeline Status
 - System Performance
 - Model Performance

Machine Learning Pipelines

Pipeline Orchestration: Bringing It Together

- ❑ Steps within **Task** can be automated
- ❑ Each set of tasks has **Inputs & Artifacts** produced as part of those steps
- ❑ **Orchestration** is required to coordinate the execution of tasks and steps within the tasks.




Model Lineage & Artifact Tracking



Model Lineage

What is Model Lineage?

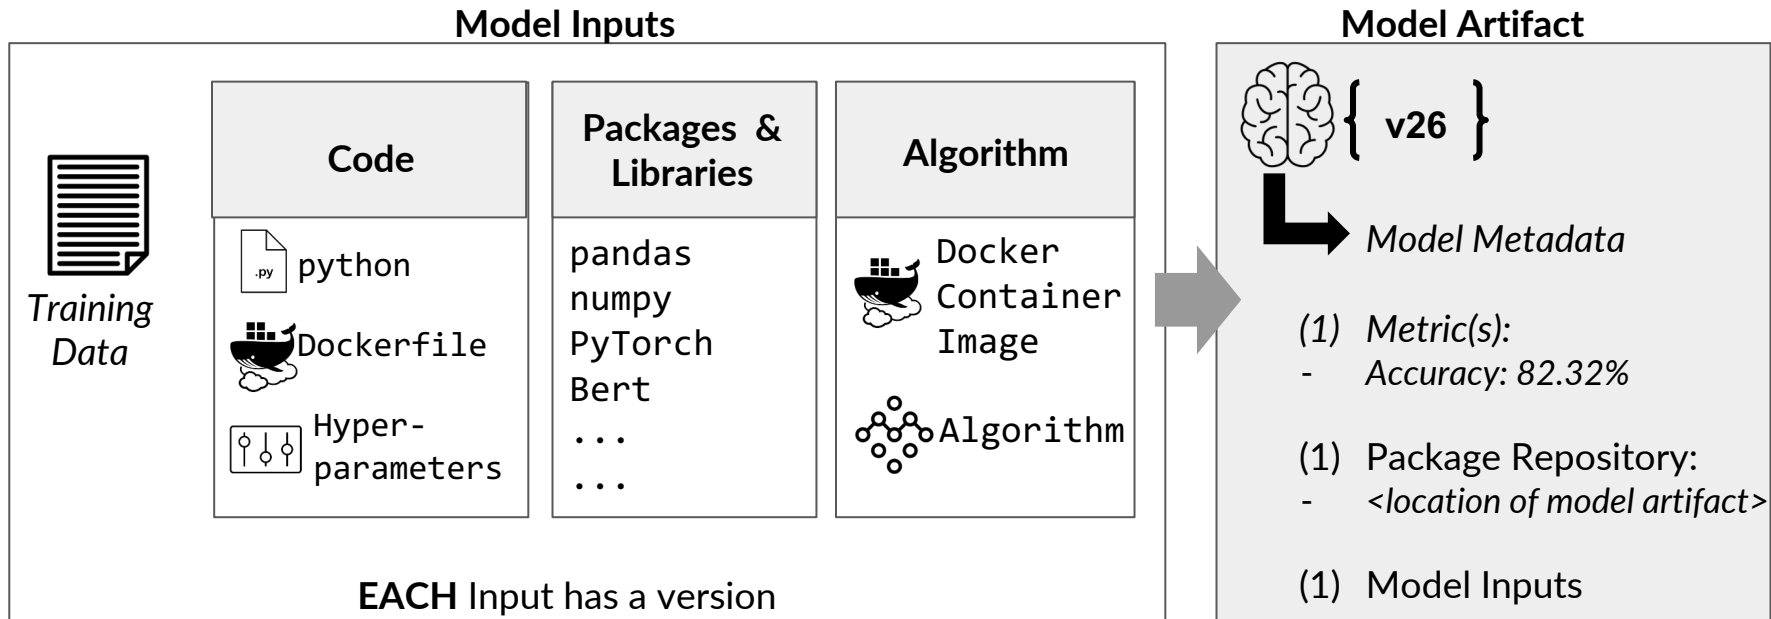
For **EACH** version of a trained model:



- ☐ Version(s) of data used
- ☐ Version(s) of code/hyperparameters used
- ☐ Version(s) of algorithm/framework
- ☐ Version(s) of training docker image
- ☐ Version(s) of packages/libraries

Model Lineage


Model Lineage Example



Model Lineage

Model Registry


What is a Model Registry?



- ❑ Centrally manage model metadata and model artifacts
- ❑ Track which models are deployed across environments

Artifact Tracking

What is Artifact Tracking?



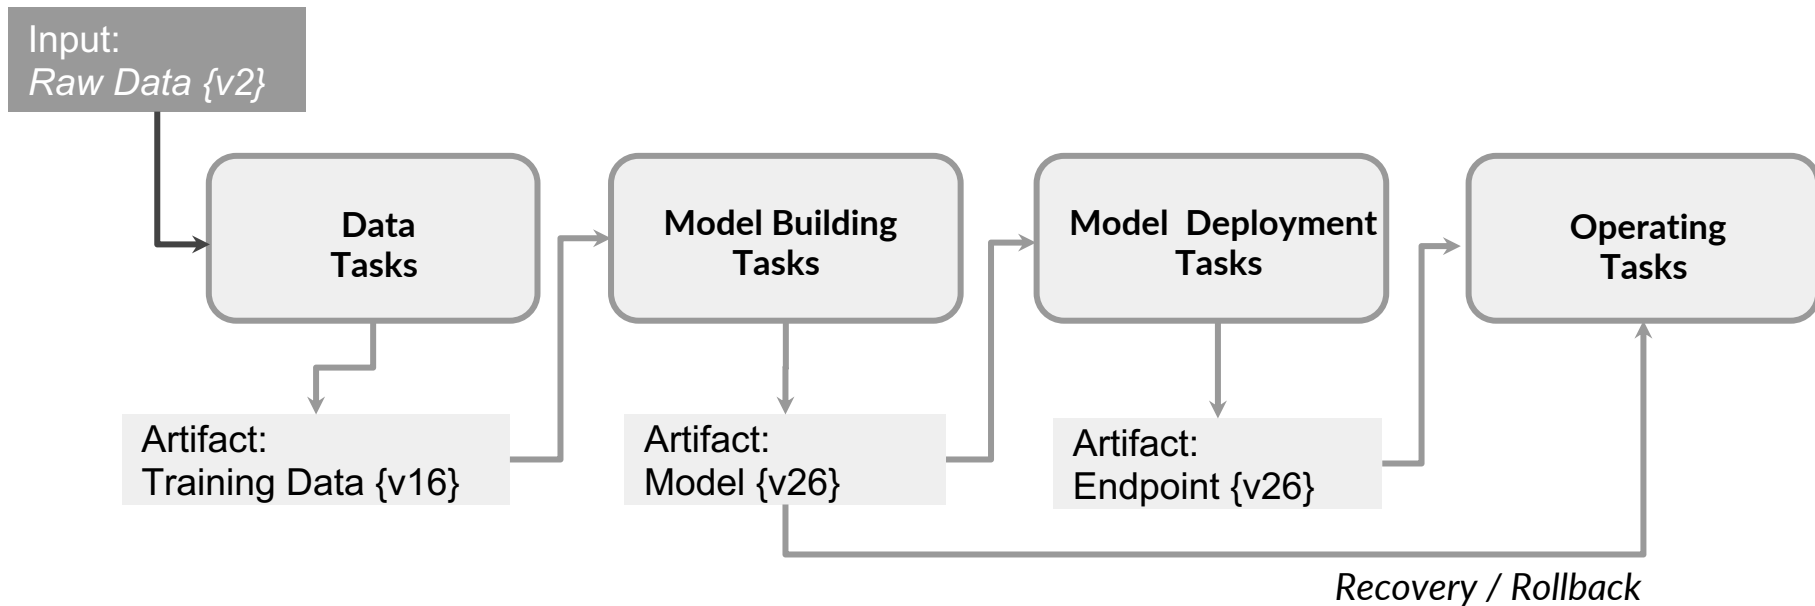
- An **Artifact** is the output of a step or task can be consumed by the next step in a pipeline or deployed directly for consumption



Artifact Tracking

Pipeline Manifest

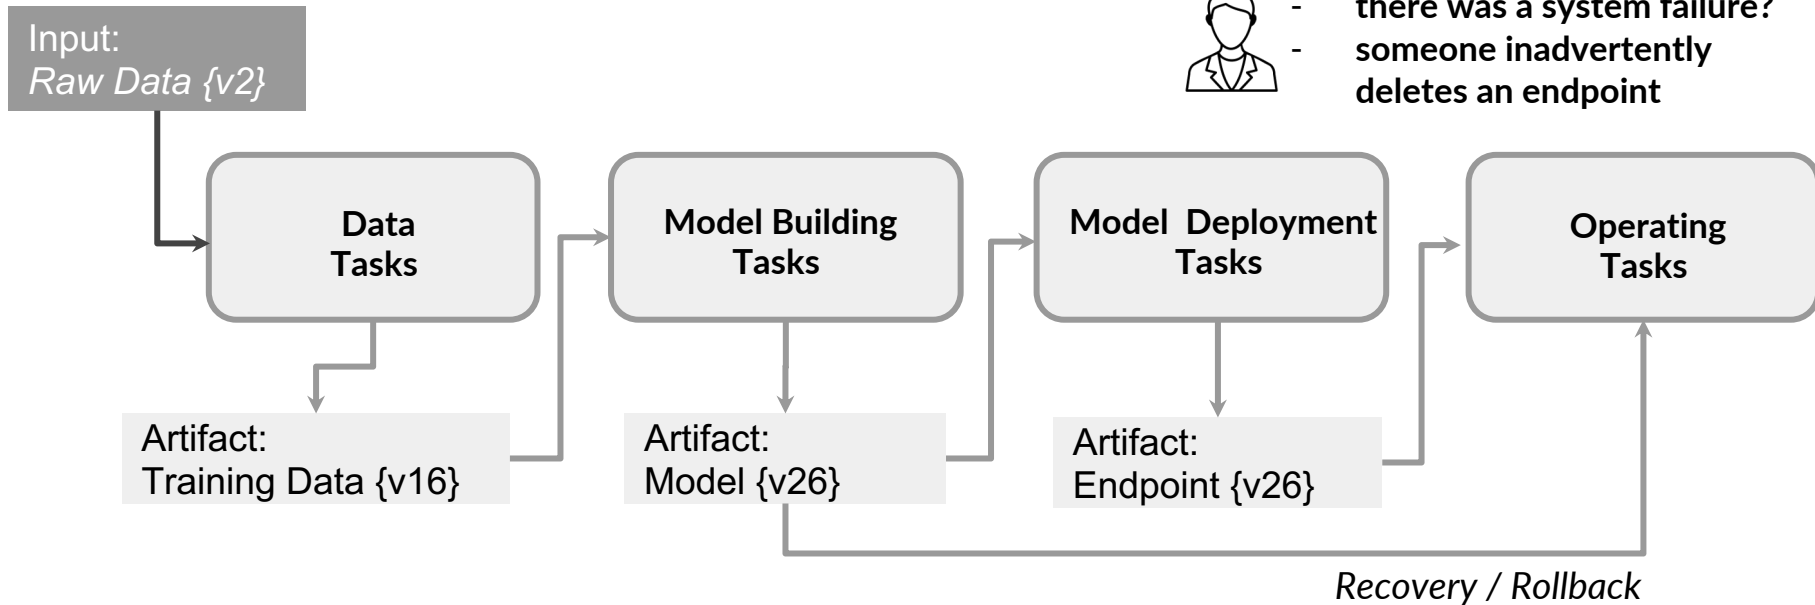
Example:



Artifact Tracking

Pipeline Manifest - Why it matters

Example:

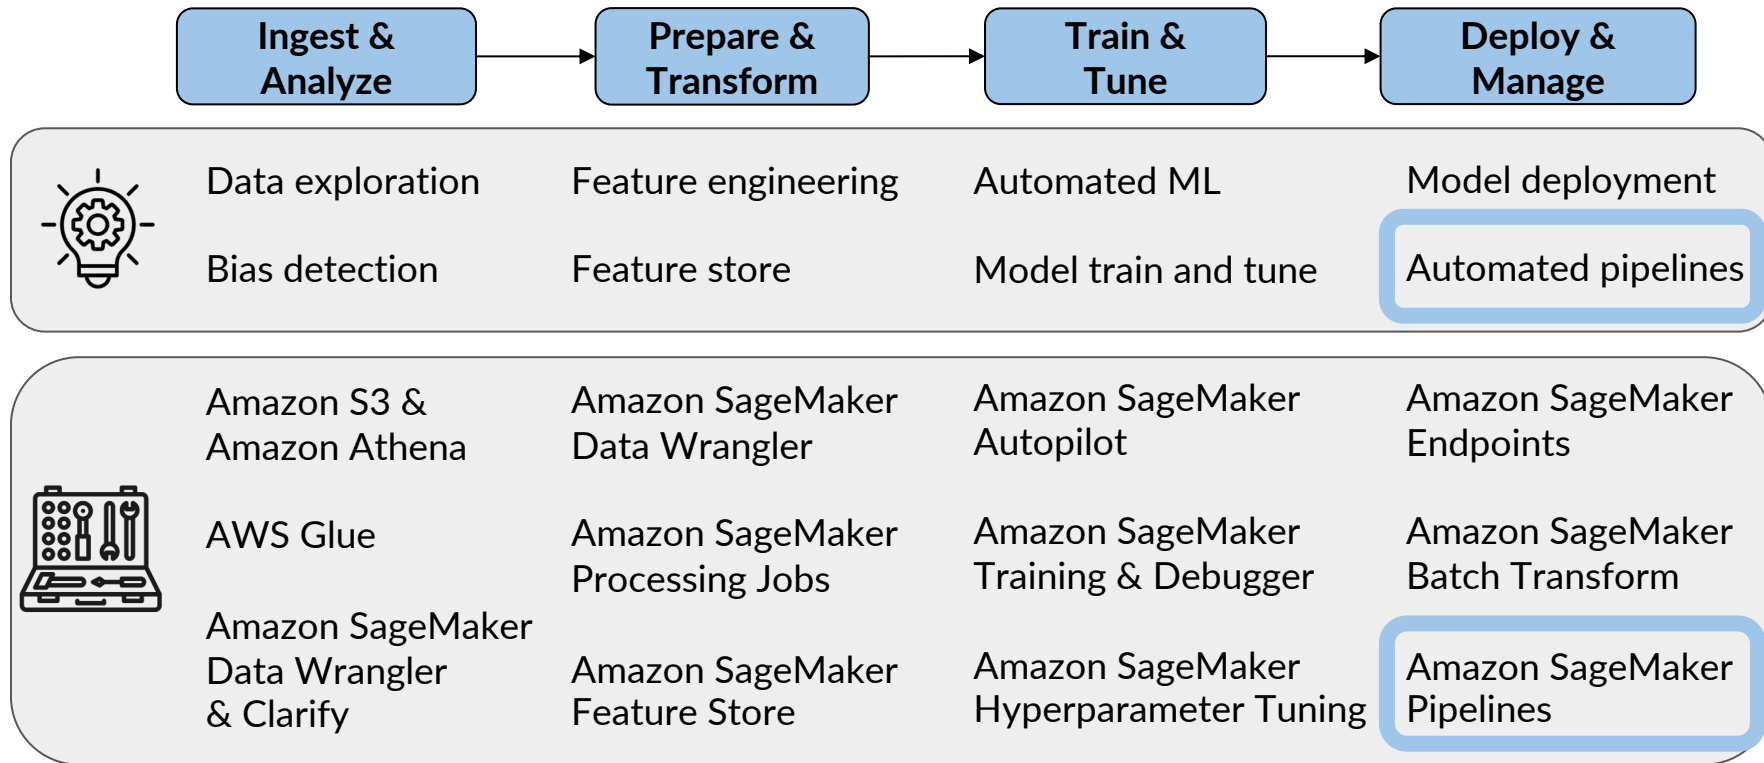


Machine Learning Pipelines


with
Amazon SageMaker Pipelines



Machine Learning Workflow



Amazon SageMaker Pipelines




Create & visualize
automated workflows



Choose the best
performing model
to deploy




Automatic
tracking of models



Bring CI/CD to
Machine Learning

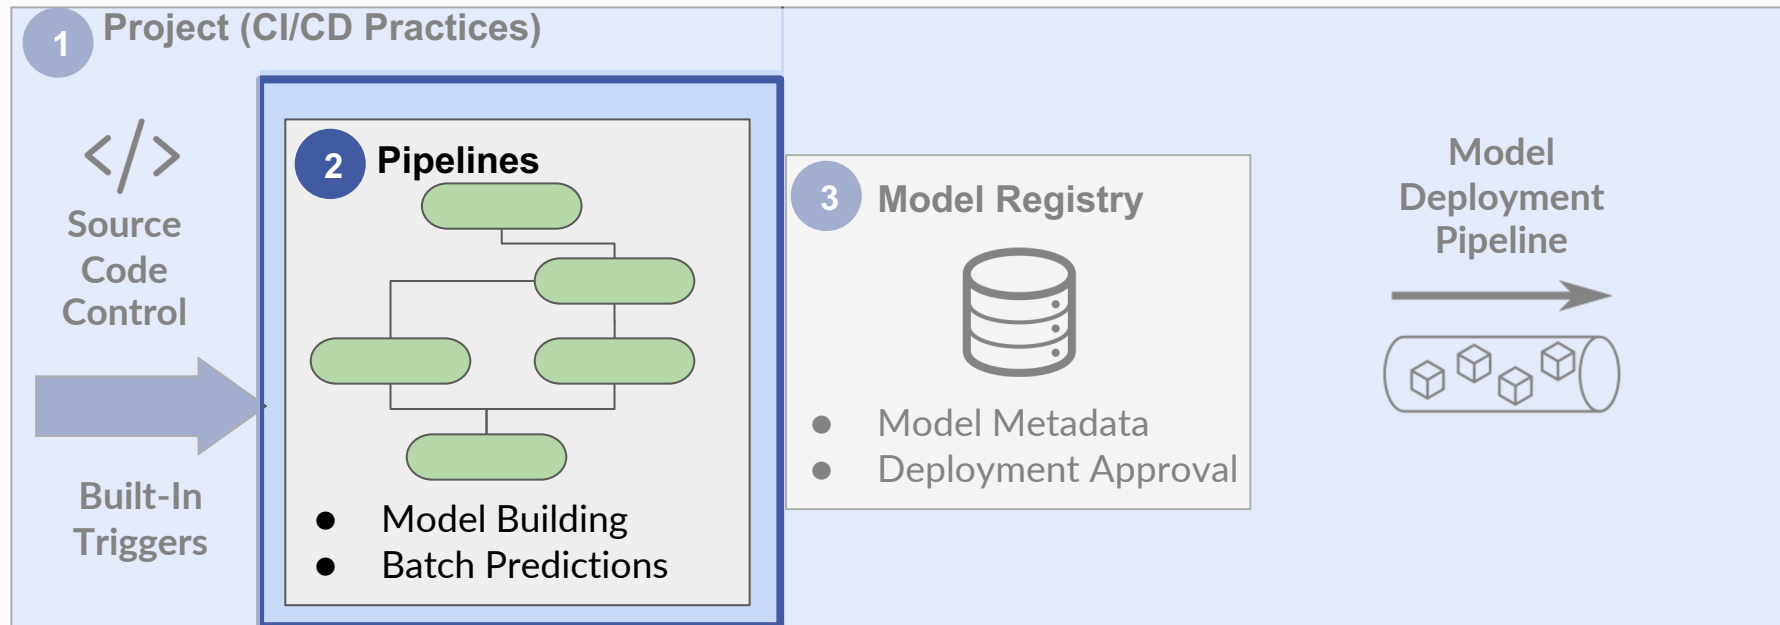
Amazon SageMaker Pipelines

SageMaker Pipelines has 3 components



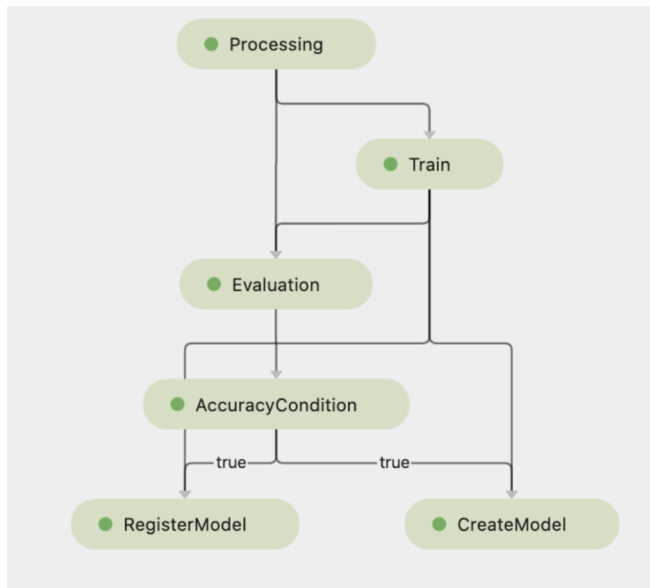
Amazon SageMaker Pipelines

SageMaker Pipelines has 3 components



SageMaker Pipelines

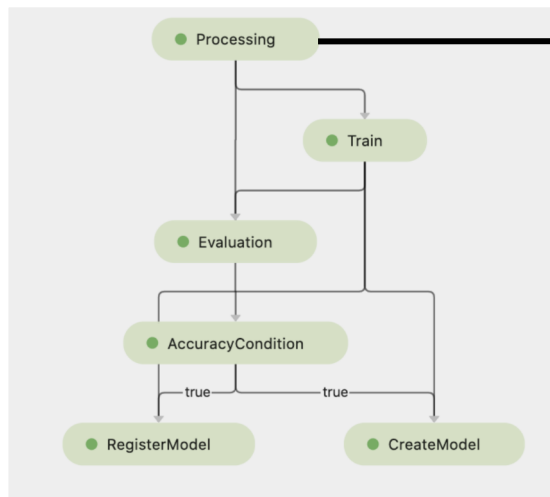
Pipelines



- ❑ Create Pipelines to build and evaluate models
- ❑ Python SDK for building workflows
- ❑ Pipeline visualization available through Amazon SageMaker Studio
- ❑ Fully managed pipelines - no servers to manage

SageMaker Pipelines

Processing Step




Use Amazon SageMaker Processing to process data set for training →





S3

- 1) Product Reviews Dataset
- 2) Processing Script



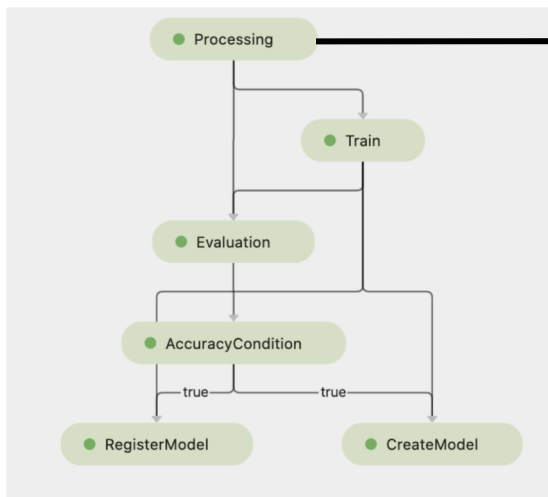
S3

- 1) Train/Validation/Test Datasets



SageMaker Pipelines

Processing Step




Define Step Inputs & Outputs →

```
processing_inputs = [  
    ProcessingInput(  
        input_name='customer-reviews-input-data',  
        source='s3://...',  
        destination='/opt/ml/processing/input/data/',  
        s3_data_distribution_type='ShardedByS3Key'  
    )  
]  
  
processing_outputs=[  
    ProcessingOutput(...)  
]
```

SageMaker Pipelines

Processing Step

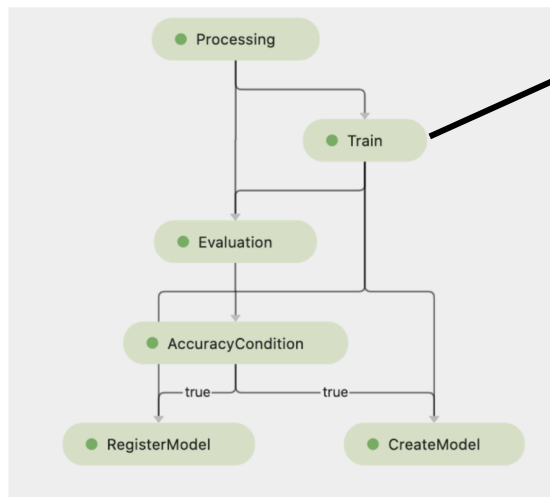


Configure the Processing Step →

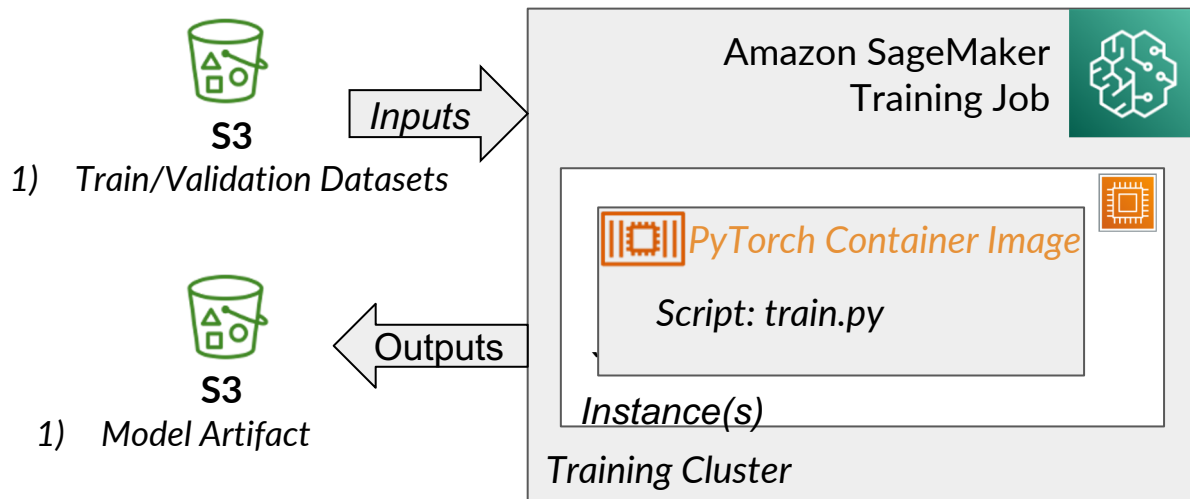
```
processing_step = ProcessingStep(  
    name='Processing',  
    code='src/prepare_data.py',  
    processor=processor,  
    inputs=processing_inputs,  
    outputs=processing_outputs,  
    job_arguments=[  
        '--train-split-percentage',  
        str(train_split_percentage.default_value,  
        ...  
    ]  
)
```

SageMaker Pipelines

Training Step




Use Amazon SageMaker Training Jobs to train the model using the outputs from the previous step as input →



SageMaker Pipelines

Training Step

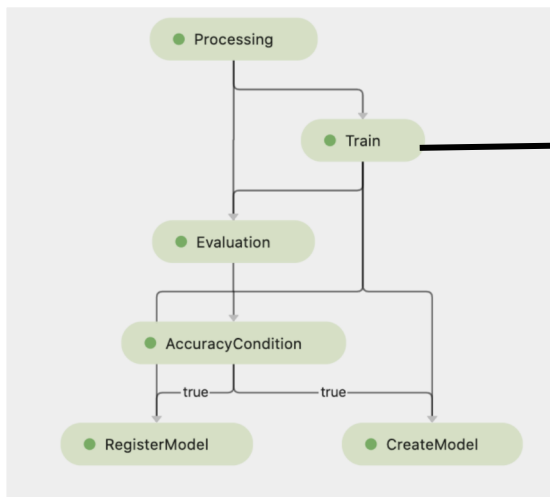


Configure Hyperparameters →

```
hyperparameters={  
    'max_seq_length': max_seq_length,  
    'epochs': epochs,  
    'learning_rate': learning_rate,  
    ...  
}
```

SageMaker Pipelines

Training Step

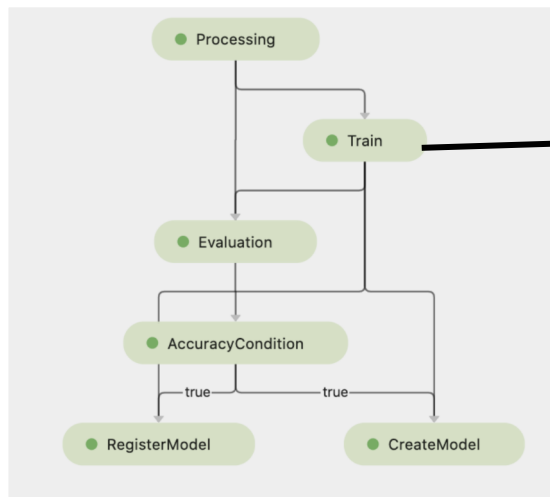


Configure Estimator →

```
from sagemaker.pytorch import PyTorch as PyTorchEstimator
estimator = PyTorchEstimator(
    entry_point='train.py',
    source_dir='src',
    role=role,
    instance_count=train_instance_count,
    instance_type=train_instance_type,
    volume_size=train_volume_size,
    py_version='py3',
    framework_version='1.6.0',
    hyperparameters=hyperparameters,
    metric_definitions=metric_definitions,
    input_mode=input_mode
)
```

SageMaker Pipelines


Training Step



Configure the Training Step →

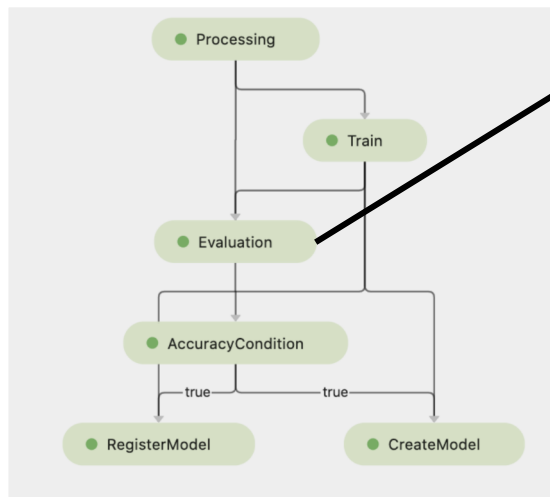
```
training_step = TrainingStep(
    name='Train',
    estimator=estimator,
    inputs={
        'train': TrainingInput(
            s3_data=processing_step.properties.ProcessingOutputConfig.Outputs[
                'sentiment-train'
            ].S3Output.S3Uri,
            content_type='text/csv'
        ),
        'validation': TrainingInput(...)
    }
)
```

Output from
Processing Step

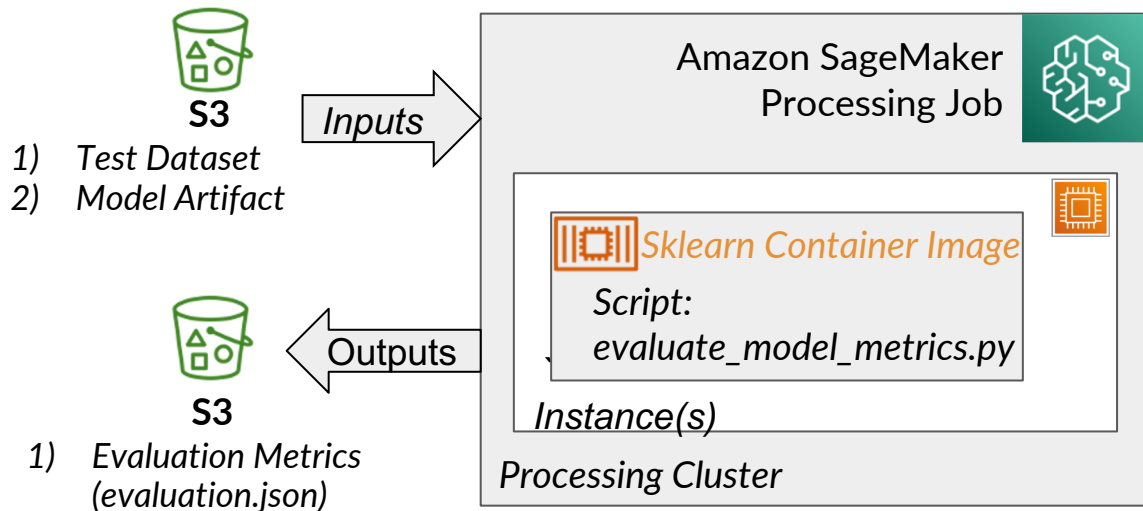


SageMaker Pipelines

Evaluation Step

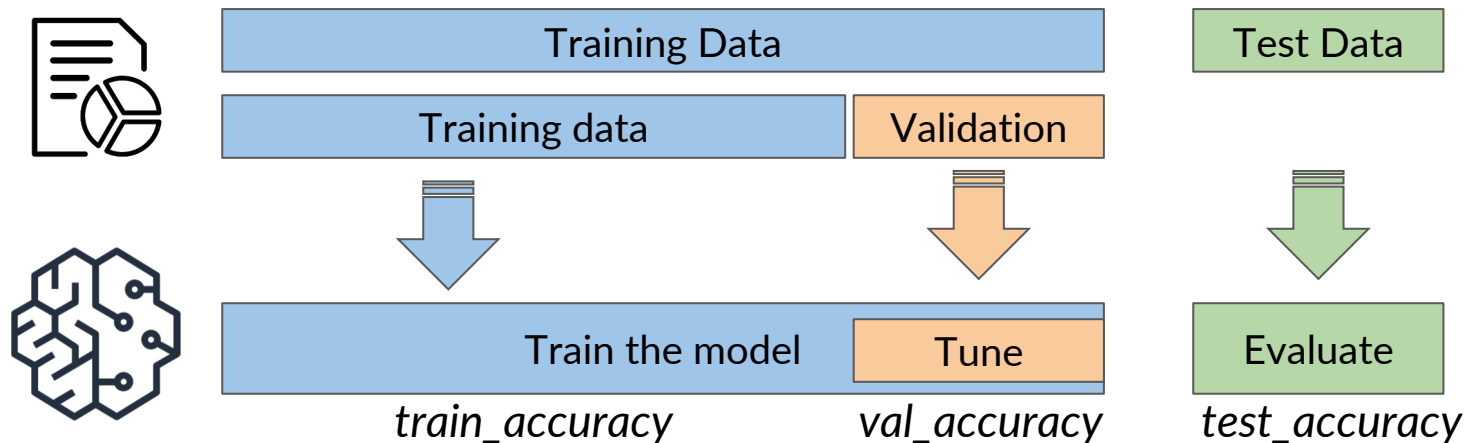


Use Amazon SageMaker Processing to evaluate trained model using test holdout dataset →



Model evaluation

- Evaluate the model with holdout test dataset



Code: *evaluate_model_metrics.py*

```
from sklearn.metrics import classification_report
from sklearn.metrics import accuracy_score
...
```

```
def predict_fn(input_data, model):
    model.eval()
    ...
    return predicted_classes_jsonlines
```

Define model
predict
function

```
...
y_test = df_test_reviews['review_body'].map(predict)
y_actual = df_test_reviews['sentiment'].astype('int64')
```

Use "test"
holdout data

```
print(classification_report(y_true=y_test, y_pred=y_actual))
```

```
accuracy = accuracy_score(y_true=y_test, y_pred=y_actual)
print('Test accuracy: ', accuracy)
```

Calculate
test accuracy

Analyze results


```
from pprint import pprint

evaluation_json = sagemaker.s3.S3Downloader.read_file(
    "{}evaluation.json".format(evaluation_metrics_s3_uri))

pprint(json.loads(evaluation_json))
```

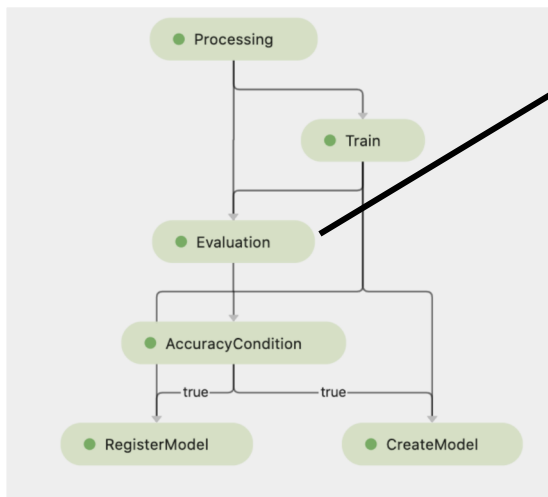


```
>> {'metrics': {'accuracy': {'value': 0.7458165031736872}}}
```



SageMaker Pipelines

Evaluation Step




Define Output →

```
from sagemaker.workflow.properties import  
PropertyFile  
  
evaluation_report = PropertyFile(  
    name='EvaluationReport',  
    output_name='metrics',  
    path='evaluation.json'  
)
```

SageMaker Pipelines

Evaluation Step

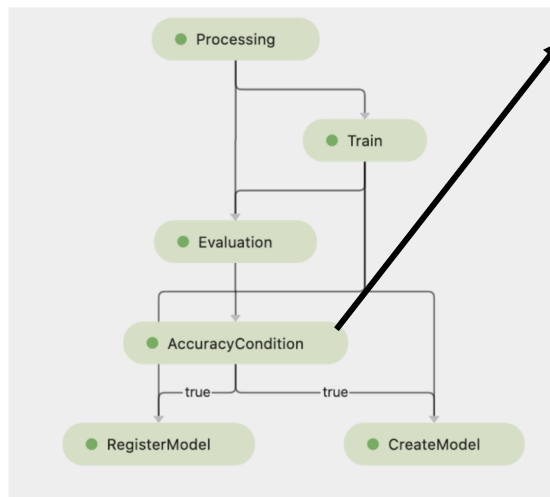


Configure Processing Step →

```
evaluation_step = ProcessingStep(  
    name='EvaluateModel',  
    processor=evaluation_processor,  
    code='src/evaluate_model_metrics.py',  
    inputs=[  
        ProcessingInput(...),...  
    ],  
    outputs=[  
        ProcessingOutput(...),  
    ],  
    job_arguments=[...],  
    property_files=[evaluation_report],  
)
```


SageMaker Pipelines

Condition Step






Use Amazon SageMaker Pipelines *Condition Step* to conditionally execute step(s) →

Evaluation Step




Conditional Step



SageMaker Pipelines

Condition Step



Define a Condition & Import Conditional Workflow Step →

```
min_accuracy_value = ParameterFloat(  
    name="MinAccuracyValue",  
    default_value=0.01  
)
```

```
from sagemaker.workflow.conditions import  
ConditionGreaterThanOrEqualTo
```

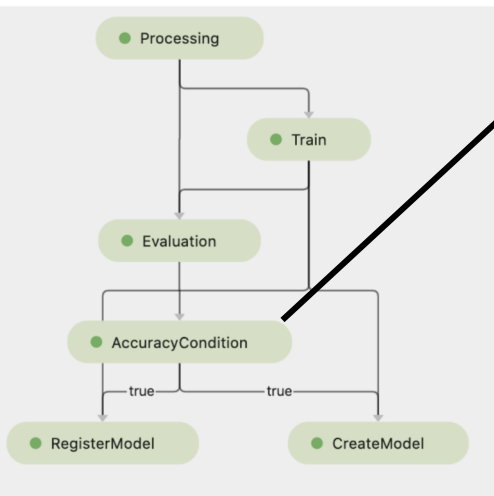
```
from sagemaker.workflow.condition_step import (  
    ConditionStep,  
    JsonGet,  
)
```


SageMaker Pipelines

Condition Step

Configure the Condition →

```
minimum_accuracy_condition =  
ConditionGreaterThanOrEqualTo(  
    left=JsonGet(  
        step=evaluation_step,  
        property_file=evaluation_report,  
        json_path="metrics.accuracy.value",  
    ),  
    right=min_accuracy_value # accuracy  
)
```




SageMaker Pipelines

Condition Step

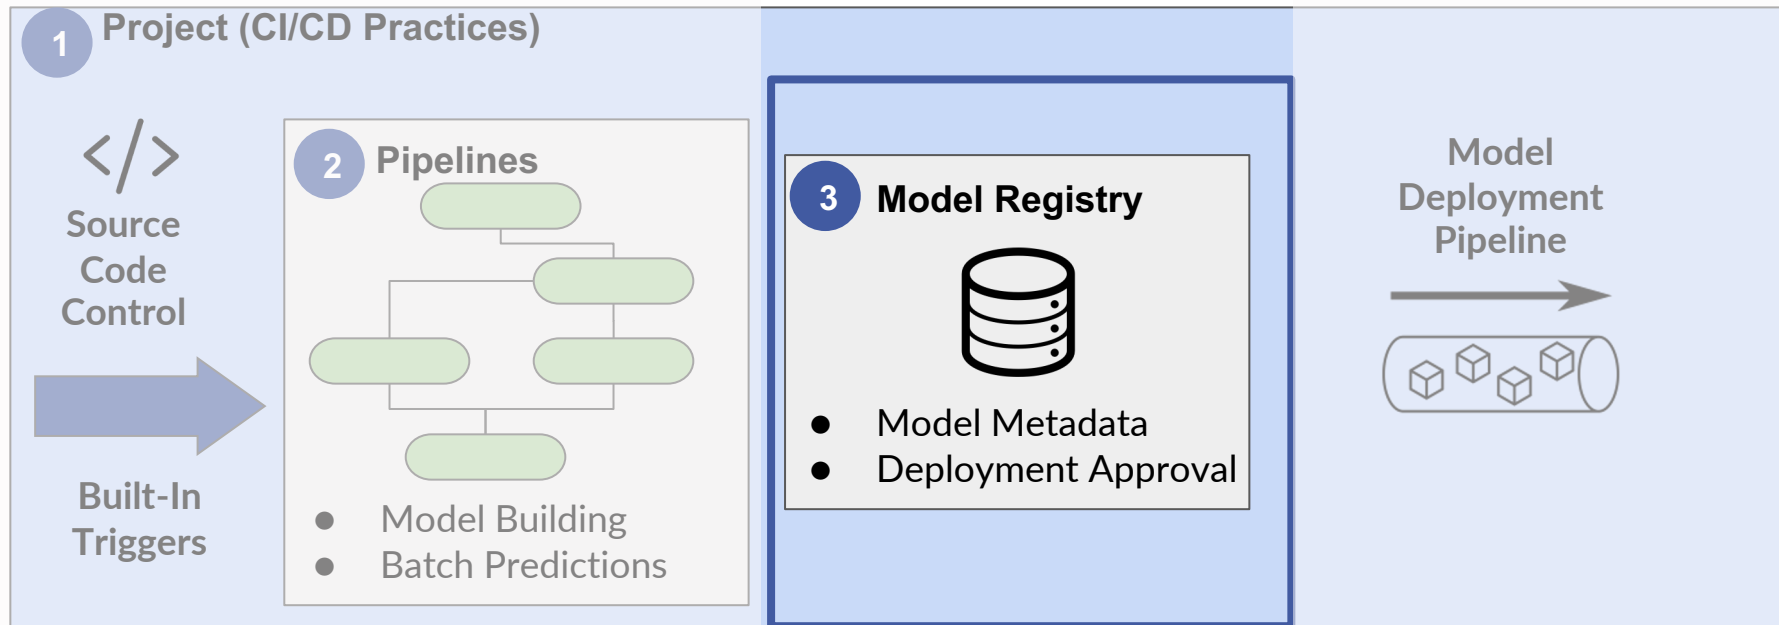
Configure the Condition Step →

```
minimum_accuracy_condition_step = ConditionStep(  
    name="AccuracyCondition",  
    conditions=[minimum_accuracy_condition],  
    # success, continue with model registration  
    if_steps=[register_step, create_step],  
    else_steps=[], # fail, end the pipeline  
)
```



Amazon SageMaker Pipelines

SageMaker Pipelines has 3 components



SageMaker Pipelines

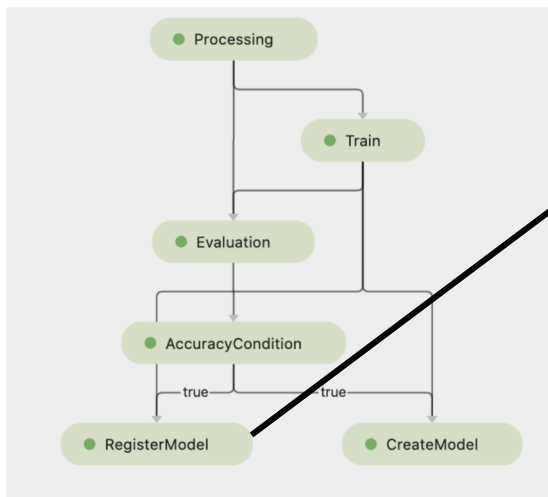
Model Registry



- ❑ Catalog models for production
- ❑ Manage model versions & metadata
- ❑ Manage the approval status of a model
- ❑ Trigger model deployment pipeline

SageMaker Pipelines

Register Model Step

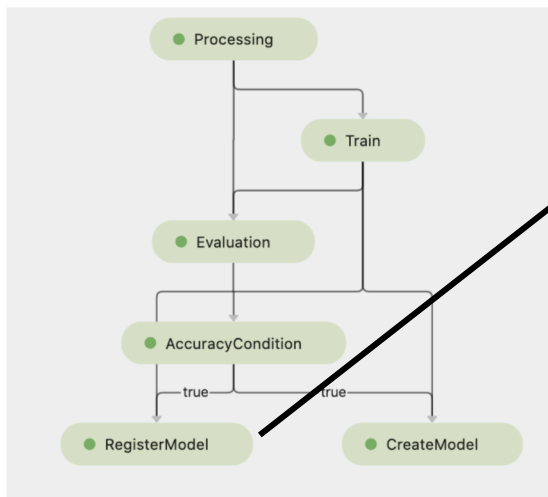


Define deployment image for inference →

```
inference_image_uri = sagemaker.image_uris.retrieve(  
    framework="pytorch",  
    region=region,  
    version="1.6.0",  
    py_version="py36",  
    instance_type=deploy_instance_type,  
    image_scope="inference"  
)
```

SageMaker Pipelines

Register Model Step



Define model metrics to be stored as metadata →

```
from sagemaker.model_metrics import MetricsSource,  
ModelMetrics
```

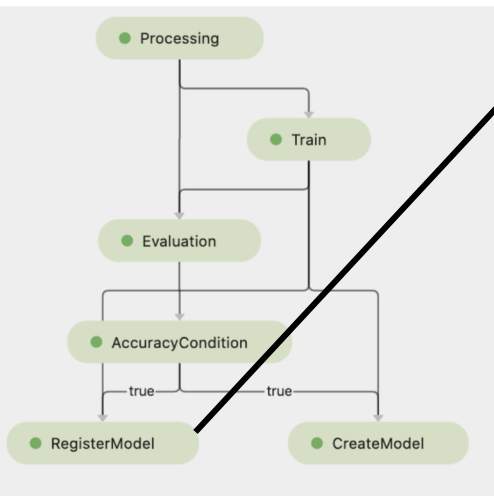
```
model_metrics = ModelMetrics(  
    model_statistics=MetricsSource(  
        s3_uri="s3://...",  
        content_type="application/json"  
    )  
)
```

SageMaker Pipelines

Register Model Step

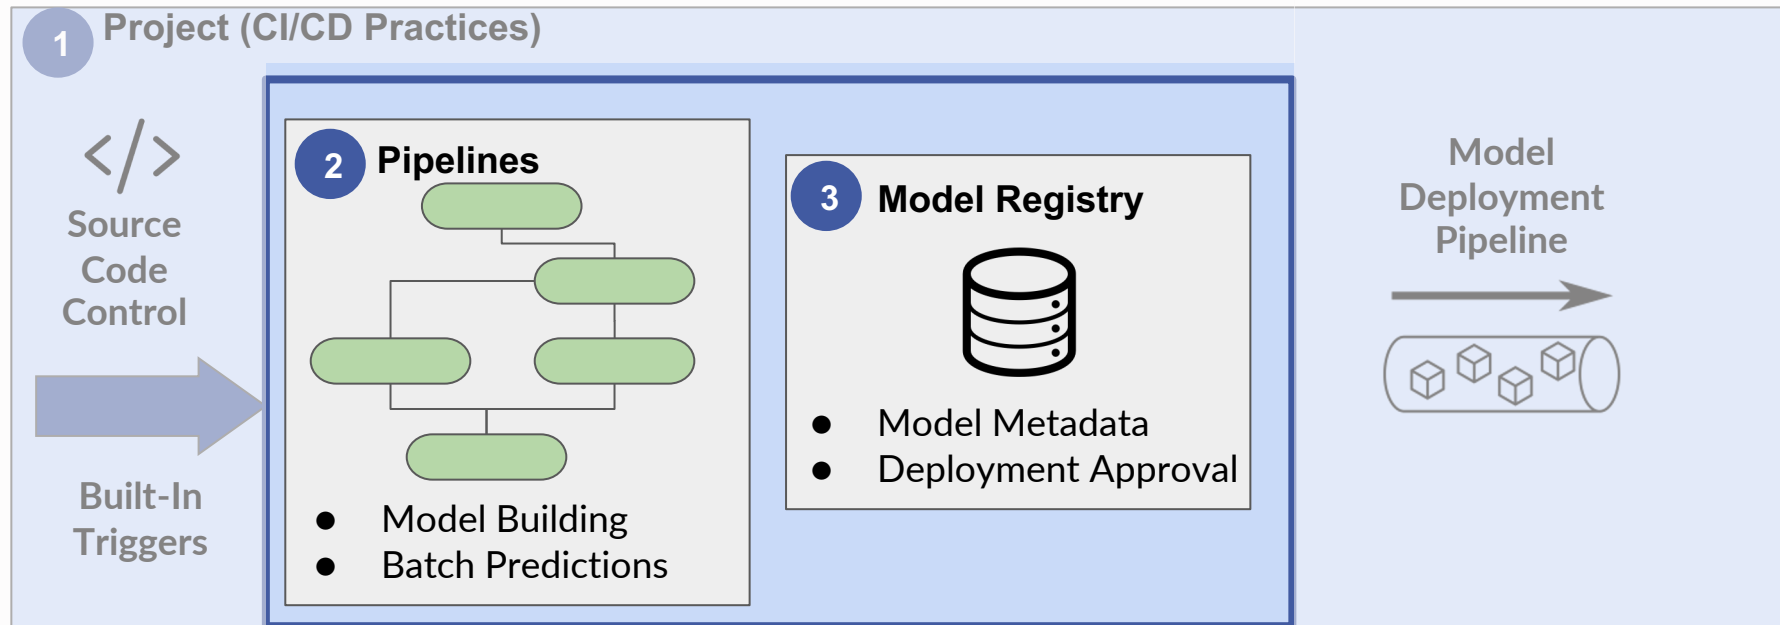
Configure the Register Model Step →

```
register_step = RegisterModel(  
    name="RegisterModel",  
    estimator=estimator,  
    image_uri=...,  
    model_data=  
        training_step.properties.ModelArtifacts.S3ModelArtifacts,  
    content_types=["application/jsonlines"],  
    response_types=["application/jsonlines"],  
    inference_instances=[deploy_instance_type],  
    transform_instances=['ml.m5.xlarge'], # batch transform  
    model_package_group_name=model_package_group_name,  
    approval_status=model_approval_status,  
    model_metrics=model_metrics)
```




Amazon SageMaker Pipelines

SageMaker Pipelines has 3 components



SageMaker Pipelines

Bringing It All Together



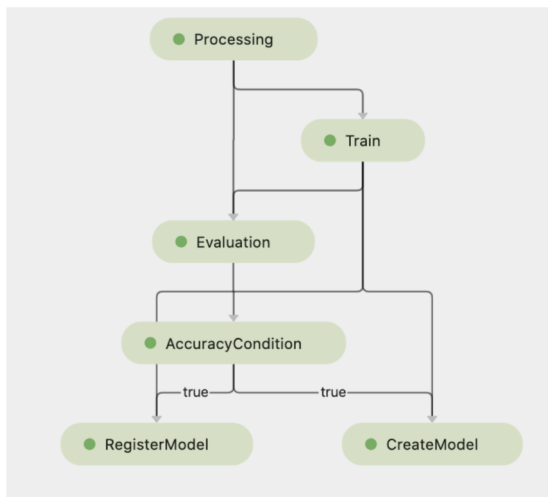
Configure the Pipeline →

```
from sagemaker.workflow.pipeline import Pipeline

pipeline = Pipeline(
    name=pipeline_name,
    parameters=[
        input_data,
        processing_instance_count,
        ...
    ],
    steps=[processing_step, training_step, evaluation_step,
           minimum_accuracy_condition_step], sagemaker_session=sess,
)
```

SageMaker Pipelines

Bringing It All Together



Create & Execute the Pipeline →

```
response = pipeline.create(role_arn=role)

pipeline_arn = response["PipelineArn"]

execution = pipeline.start(
    parameters=dict(
        InputData=raw_input_data_s3_uri,
        ProcessingInstanceCount=1,
        ...
    )
)
```

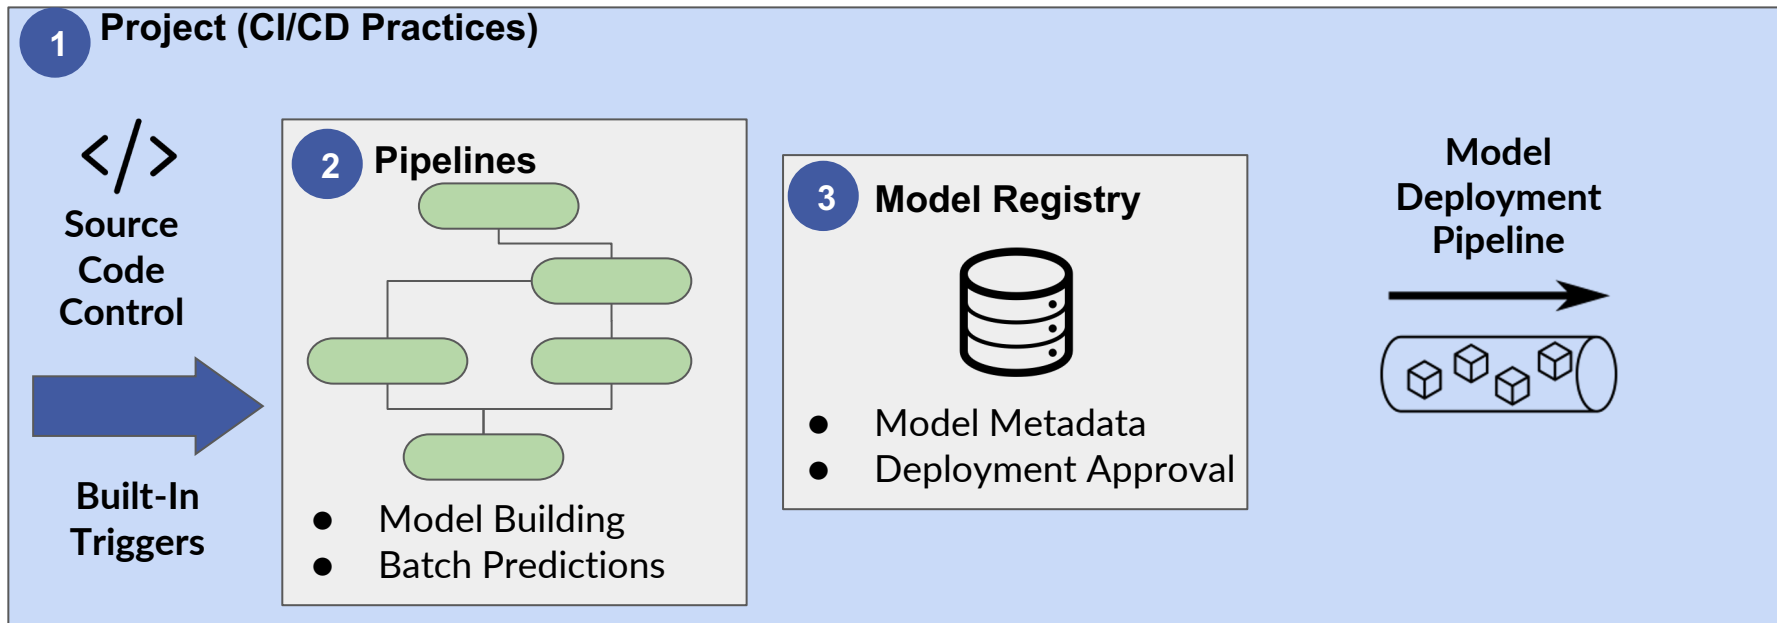
Machine Learning Pipelines

with
Amazon SageMaker Projects



Amazon SageMaker Pipelines

SageMaker Pipelines has 3 components



SageMaker Projects



- Create end-to-end ML solutions with CI/CD practices
- Incorporates source/version control
- 3 Built-In MLOps Project Templates covering:
 - Build, Train, Deploy
 - Build, Train
 - Deploy
- Ability to bring custom Project templates